

SQL 3

The Data Bootcamp | February 16, 2018

Subqueries

- A query within a query
- A subquery can almost always be replaced by a join
 - Joins may be faster.
 - SQL Server, for example, will often optimize these and rewrite them as joins under the hood. Other database implementations vary.
- Which should we do? Focus on designing a readable query first.

```
select *  
from inventory  
where film_id IN  
(  
    select film_id  
    from film  
    where title = 'Early Home'  
);
```

Subqueries

SELECT LINE

- Access information from tables in SELECT LINE
 - Print out orders and compare to average product price
- Query must return a single value

Food Truck Example:

```
sql_query = """
select avg(Price)
from Products
"""
products = pd.read_sql_query(sql_query, engine)
products.head()
```

	avg(Price)
0	6.25

```
sql_query = """
select Name, Price - (select avg(Price) from Products)
from Products;
"""
products = pd.read_sql_query(sql_query, engine)
products.head()
```

	Name	Price - (select avg(Price) from Products)
0	carne asada	5.75
1	diet coke	-3.75
2	pollo burrito	1.75
3	sprite	-3.75

Subqueries

Where Line

- Useful for comparing values from other tables
 - Find customers who have placed orders
 - Find orders containing a particular product category

We did a similar example to this one Thursday with the Sakila database using joins.

```
sql_query = """
select *
from inventory
where film_id IN
(
    select film_id
    from film
    where title = 'Early Home'
);
"""
joined = pd.read_sql_query(sql_query, engine)
joined.head()
```

	inventory_id	film_id	store_id	last_update
0	1206	268	2	2006-02-15 05:09:17
1	1207	268	2	2006-02-15 05:09:17

OLTP, OLAP, ETL, ELT, Oh my...

Sr Data Engineer

American National Bank ★★★★★ 58 reviews - Omaha, NE

Apply Now

Experience, Education & Training Requirements

Experience Needed:

- Requires 7+ years of related experience in Enterprise Data Engineering/Management
- Experience developing and maintaining an enterprise wide data governance framework
- Experience with big data technologies
- Experience using data management software and BI Reporting tools like Cognos
- Experience writing and maintaining SQL
- Experience with master data management, data modeling and database design, OLTP and OLAP, data warehousing ETL/ELT workflows
- Financial services/banking industry experience is desired

Data Warehouse Architect

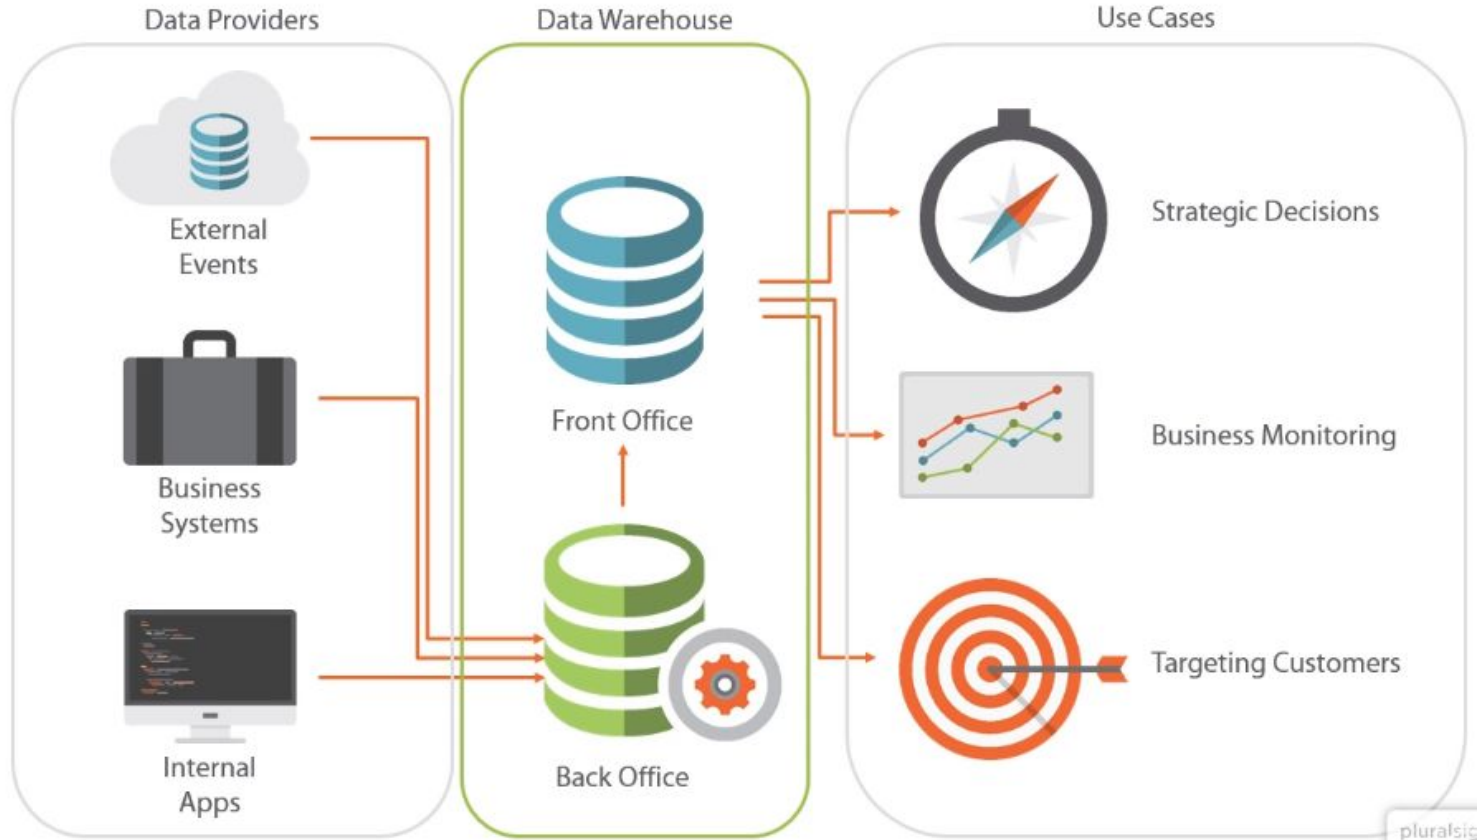
ClickBank ★★★★★ 11 reviews - Broomfield, CO

Apply Now

Key Responsibilities include, but are not limited to:


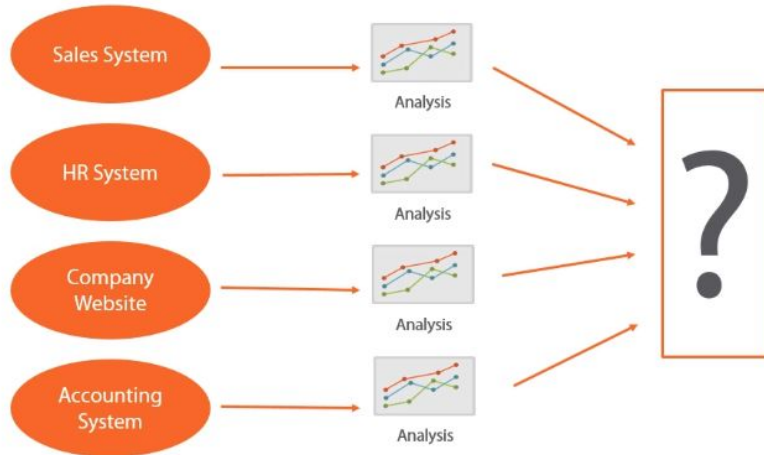
- Works with the business intelligence and the ETL development team to extract data from the OLTP system and present it in a format that supports data mining, analysis and reporting.
- Ability to maintain existing and develop new ETL jobs which support extracting source systems, transforming data based on business rules and requirements and loading of data into warehouse structures.
- Troubleshoot data issues and implement fixes to ensure data accuracy.
- Good communication skills with both technical and non-technical staff members.
- Able to translate business requirements and design ETL jobs to support them.
- Ability to re-design data models to support data mining and enhanced analytics efforts.

The Big Picture

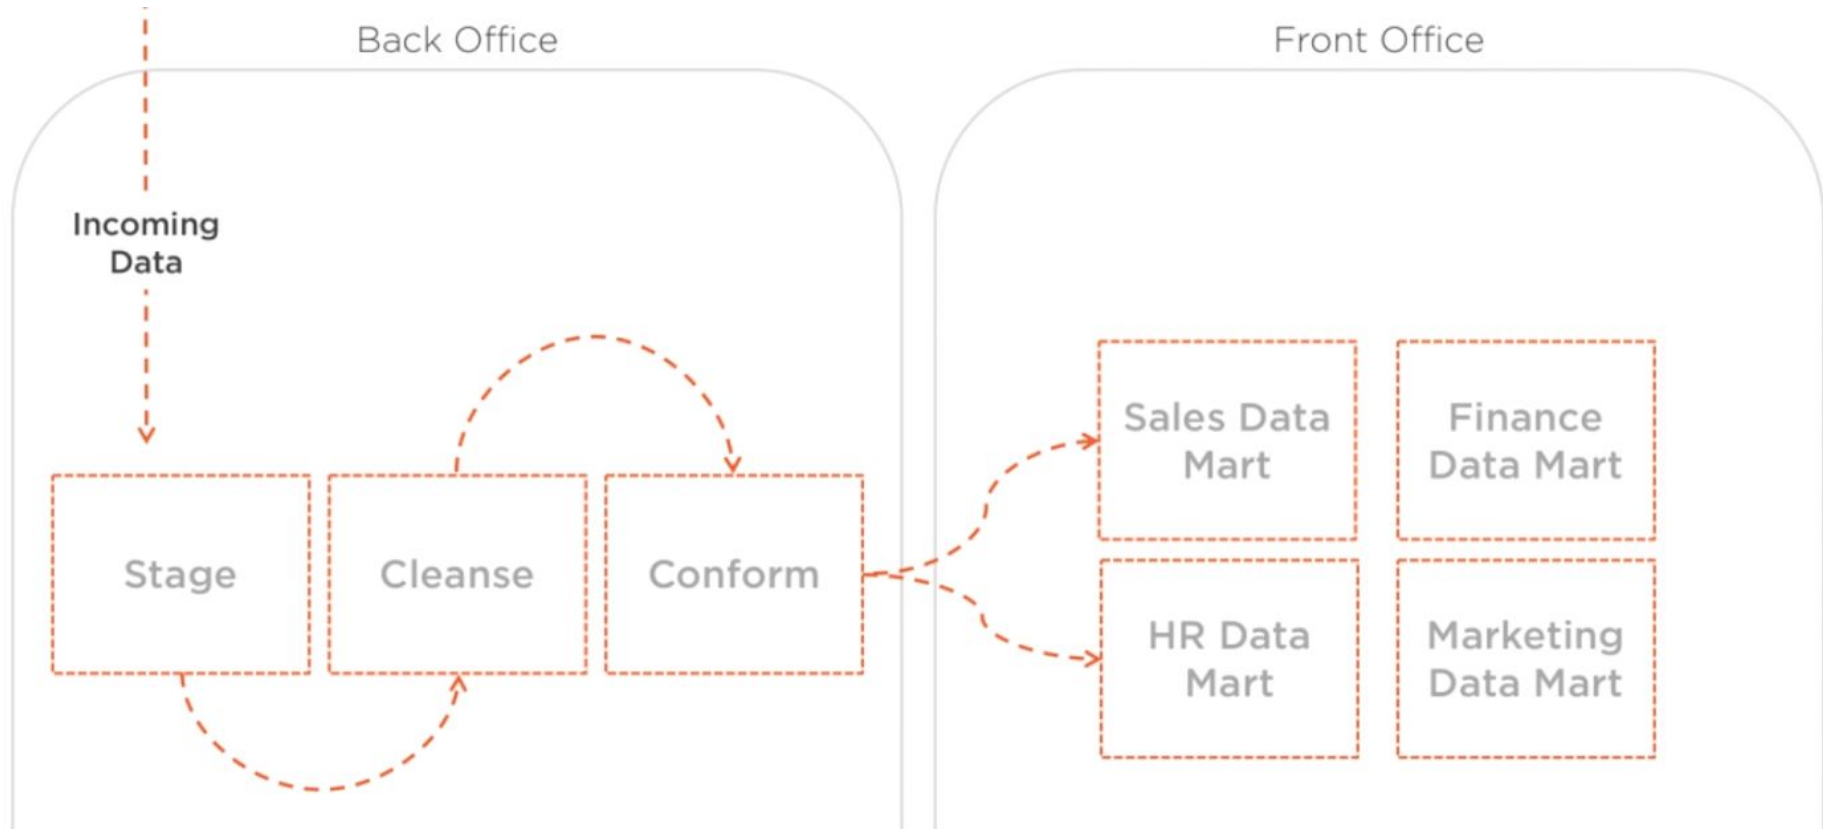


Data Warehousing

Data warehousing opens up opportunities to gain value from your raw data through the process of analysis.



The Big Picture



The Big Picture

Front Office

Sales Data Mart

Finance Data Mart

HR Data Mart

Marketing Data Mart

Front End



File Exports



BI/Analytics Tools



SQL Queries



OLAP Cubes



- Strategic Analysis
- Data Science
- Operational Efficiency
- Research & Development
- Sales Optimization
- Etc...



The Big Picture



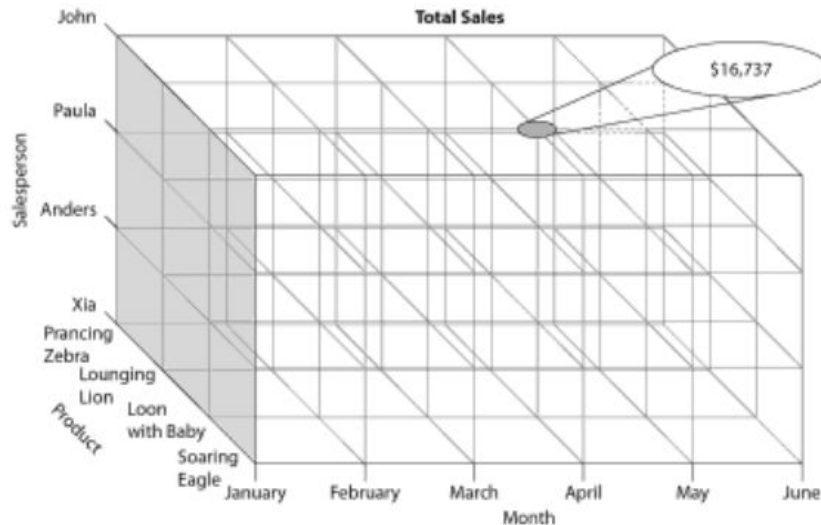
OLAP Cubes

OLAP We will go over in the next slide.

A **Cube** is a structure that contains a value for one or more measures for each unique combination of the members of all its dimensions.

OLAP Cubes are fading as Tableau / PowerBI / Pandas are moving into this space.

That said, they are still out there.. everywhere...



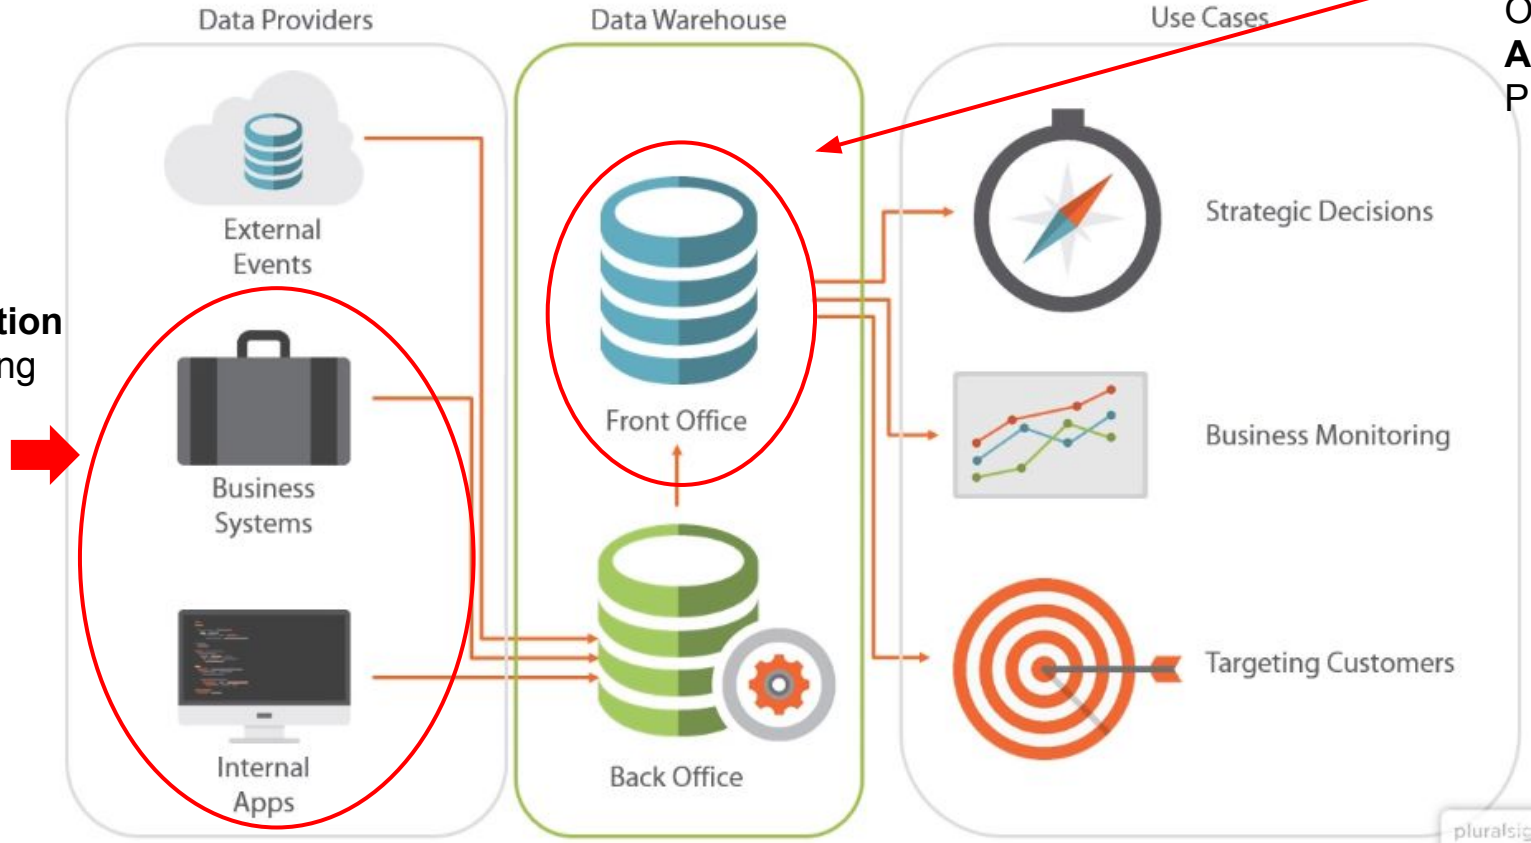
More Information. **Delivering Business Intelligence with Microsoft SQL Server 2016**

<https://www.amazon.com/exec/obidos/ASIN/1259641481/alvinashcraft-20>

OLTP Versus OLAP

OLTP
Online
Transaction
Processing

OLAP
Online
Analytical
Processing



pluralsig

OLTP and OLAP

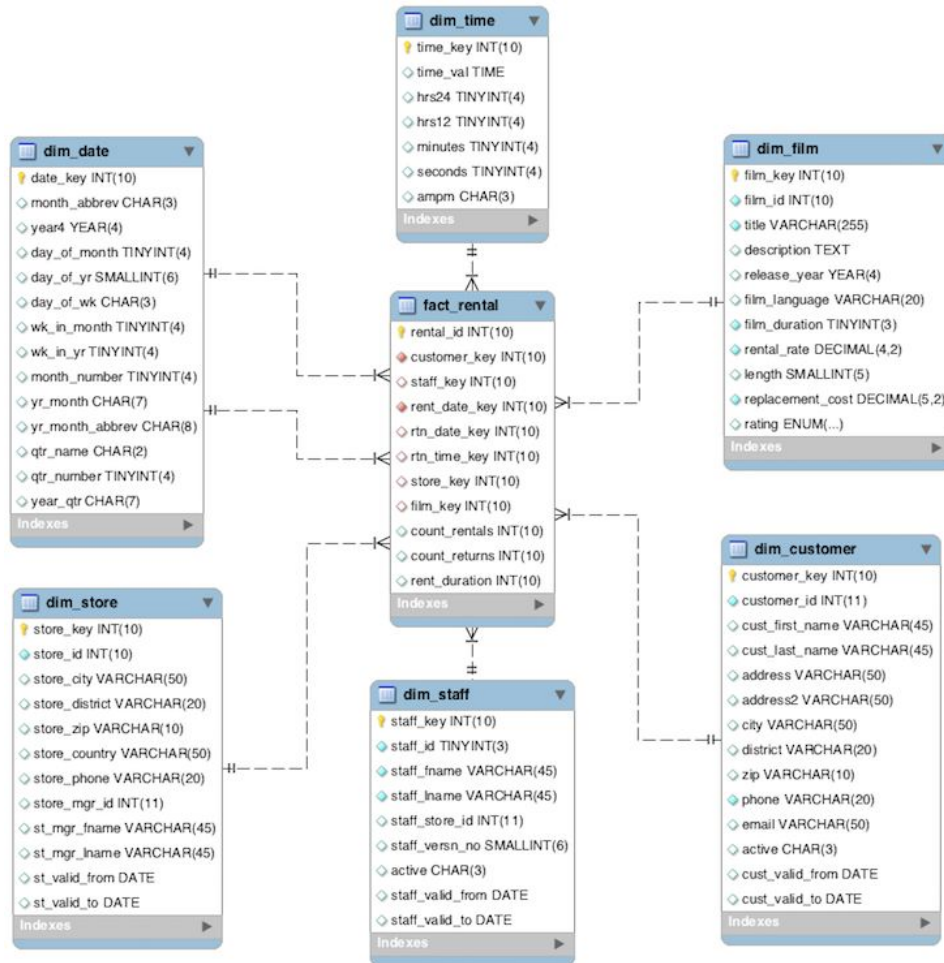
OLTP - Transactions	OLAP - Analytics
Daily operations. Always on	Periodically updated
All kinds of users	Managers, Data Analysts
Directly supports business processes. Orders, shipping, customer info	For long-term planning, not tracking orders.
Current Status. Real Time Data.	Maintains History
Entity-Relationship Model	Dimensional Model
Normalized to 3NF	Denormalized dimensions

Why have two systems?

Normalized (OLTP)	Denormalized (OLAP)
Machine oriented. Quick writes.	Human Oriented
Complex Joins	Simple Joins
Eliminates redundancy. Product name not repeated so you don't mess it up.	Use redundancy to simplify

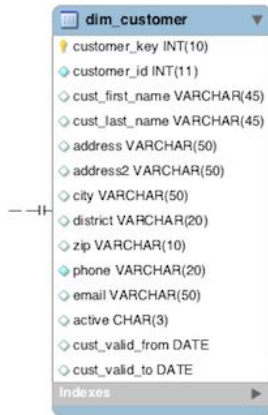
Star Schema

A data mart design pattern consisting of fact and dimension tables resembling a star like shape.

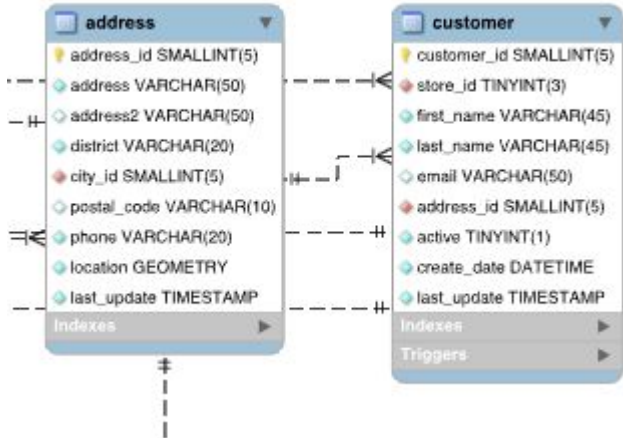


Star Schema / Normalized Schema Comparison

Star Schema - OLAP



3NF - OLTP



ETL Software



Oracle Data Integration

Next Generation Data Integration for the Enterprise



In a startup, you might do this with Python and SQLAlchemy / MongoDB.

A medium to large company at scale will almost undoubtedly use one of these ETL Software Packages



ETL Stages



Stage

Extract: Bring everything in from everywhere.



Cleanse

Transform: Cleanup names, move things around. Most of the work.



Deliver

Load: Deliver to front office database

To Learn More

Star Schema Foundations:

<https://app.pluralsight.com/library/courses/star-schema-foundations/table-of-contents>

The Kimball Group Reader

<https://www.amazon.com/gp/product/1119216311>